


---

# ATP Player Hand Analysis And Prediction

CAPSTONE PROJECT 1





# Overview

- Problem
- The data
- Data Wrangling and preparation
- Exploratory Data Analysis
- Statistical Inference
- Machine Learning



# Problem

- Do ATP left hand players have advantage against ATP right hand players?
- Can we predict if an ATP player is righty or lefty based on their match results?

# The Data

- The data used for this project comes from Github.
- The dataset includes ATP match results from 2000 to 2019

```
['data/tennis_atp/match_00_19/atp_matches_2000.csv',  
'data/tennis_atp/match_00_19/atp_matches_2001.csv',  
'data/tennis_atp/match_00_19/atp_matches_2002.csv',  
'data/tennis_atp/match_00_19/atp_matches_2003.csv',  
'data/tennis_atp/match_00_19/atp_matches_2004.csv',  
'data/tennis_atp/match_00_19/atp_matches_2005.csv',  
'data/tennis_atp/match_00_19/atp_matches_2006.csv',  
'data/tennis_atp/match_00_19/atp_matches_2007.csv',  
'data/tennis_atp/match_00_19/atp_matches_2008.csv',  
'data/tennis_atp/match_00_19/atp_matches_2009.csv',  
'data/tennis_atp/match_00_19/atp_matches_2010.csv',  
'data/tennis_atp/match_00_19/atp_matches_2011.csv',  
'data/tennis_atp/match_00_19/atp_matches_2012.csv',  
'data/tennis_atp/match_00_19/atp_matches_2013.csv',  
'data/tennis_atp/match_00_19/atp_matches_2014.csv',  
'data/tennis_atp/match_00_19/atp_matches_2015.csv',  
'data/tennis_atp/match_00_19/atp_matches_2016.csv',  
'data/tennis_atp/match_00_19/atp_matches_2017.csv',  
'data/tennis_atp/match_00_19/atp_matches_2018.csv',  
'data/tennis_atp/match_00_19/atp_matches_2019.csv']
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 61560 entries, 0 to 2780  
Data columns (total 49 columns):  
tourney_id      61560 non-null object  
tourney_name    61560 non-null object  
surface         61442 non-null object  
draw_size      2781 non-null float64  
tourney_level   61560 non-null object  
tourney_date    61560 non-null int64  
match_num       61560 non-null int64  
winner_id       61560 non-null int64  
winner_seed     25567 non-null object  
winner_entry    7346 non-null object  
winner_name     61560 non-null object  
winner_hand     61542 non-null object  
winner_ht       56229 non-null float64  
winner_ioc      61560 non-null object  
winner_age      61545 non-null float64  
loser_id        61560 non-null int64  
loser_seed      13973 non-null object  
loser_entry     12107 non-null object  
loser_name      61560 non-null object  
loser_hand      61514 non-null object  
loser_ht        53389 non-null float64  
loser_ioc       61560 non-null object  
loser_age       61529 non-null float64  
score           61559 non-null object  
best_of         61560 non-null int64  
round           61560 non-null object  
minutes         54478 non-null float64  
w_ace           55781 non-null float64  
w_df            55781 non-null float64  
w_svpt          55781 non-null float64  
w_1stIn         55781 non-null float64  
w_1stWon        55781 non-null float64  
w_2ndWon        55781 non-null float64  
w_SvGms         55781 non-null float64  
w_bpSaved       55781 non-null float64  
w_bpFaced       55781 non-null float64  
l_ace           55781 non-null float64  
l_df            55781 non-null float64  
l_svpt          55781 non-null float64  
l_1stIn         55781 non-null float64  
l_1stWon        55781 non-null float64  
l_2ndWon        55781 non-null float64  
l_SvGms         55781 non-null float64  
l_bpSaved       55781 non-null float64  
l_bpFaced       55781 non-null float64  
winner_rank     61057 non-null float64  
winner_rank_points 61057 non-null float64  
loser_rank      60263 non-null float64  
loser_rank_points 60263 non-null float64  
dtypes: float64(28), int64(5), object(16)  
memory usage: 23.5+ MB
```

# Data Wrangling and Preparation

- Removing extraneous data – i.e.,
  - Removing rows with 100% missing metric values
  - Removing some categorical columns not needed
- Feature engineering—i.e.,
  - computing and adding columns for data exploration and machine learning
  - Prepared data set of ATP matches between righties and lefties

```
df.loc[df.player_hand=='R', 'player_hand_flag'] = 1
df.loc[df.player_hand=='L', 'player_hand_flag'] = 0
```




```
df_all['svpt_won_pct'] = np.around((df_all.sv1stWon + df_all.sv2ndWon)/df_all.svpt,2)
df_all['svpt_std_var'] = df_all.svpt - df_all.SvGms*4
df_all['bpSaved_pct'] = np.around(df_all.bpSaved/df_all.bpFaced,2)
df_all.bpSaved_pct.fillna(0, inplace=True)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20132 entries, 0 to 20131
Data columns (total 48 columns):
Unnamed: 0      20132 non-null int64
tourney_id      20132 non-null object
tourney_name    20132 non-null object
surface         20132 non-null object
tourney_level   20132 non-null object
tourney_date    20132 non-null int64
match_num      20132 non-null int64
player_id       20132 non-null int64
player_name     20132 non-null object
player_hand     20132 non-null object
player_ht       20132 non-null float64
player_ioc      20132 non-null object
player_age      20132 non-null float64
score           20132 non-null object
best_of         20132 non-null int64
round           20132 non-null object
minutes         20132 non-null float64
ace             20132 non-null float64
df              20132 non-null float64
svpt            20132 non-null float64
sv1stIn         20132 non-null float64
sv1stWon        20132 non-null float64
sv2ndWon        20132 non-null float64
SvGms           20132 non-null float64
bpSaved         20132 non-null float64
bpFaced         20132 non-null float64
player_rank     20132 non-null float64
player_rank_points 20132 non-null float64
ace_pct         20132 non-null float64
df_pct          20132 non-null float64
sv1stIn_pct     20132 non-null float64
sv2ndIn_pct     20132 non-null float64
sv1stWon_pct    20132 non-null float64
sv2ndWon_pct    20132 non-null float64
GmsWon          20132 non-null float64
GmsLoss         20132 non-null float64
year            20132 non-null int64
opponent_id     20132 non-null int64
opponent_name   20132 non-null object
won_flag        20132 non-null int64
player_age_bucket 20132 non-null object
player_hand_flag 20132 non-null float64
surface_id      20132 non-null float64
tourney_level_id 20132 non-null float64
player_rank_group 20132 non-null float64
svpt_won_pct    20132 non-null float64
svpt_std_var    20132 non-null float64
bpSaved_pct     20132 non-null float64
dtypes: float64(29), int64(8), object(11)
memory usage: 7.4+ MB
```

# Exploratory Data Analysis

Some observations:



- Average height of lefties are taller than the righties played in the same year except year 2001, 2008 and 2009
- Except 2007 - 2009, the average age of lefties are older than righties
- The average of rankings of righties are higher than lefties



# Exploratory Data Analysis

Some observations:

- Compared 20 year mean ace percentage for righties and lefties, the righties had higher ace percentage than lefties
- Lefties has a little higher double fault percentage than righties
- Compared 20 year matches, lefties do have higher first in server percentage
- The first serve won percentage seems no significant difference between righties and lefties
- Righties also have higher second serve in percentage than lefties





# Statistical Inference

Some observations:

- Performed t-test to check if there are significant difference on serve stats (ace%, 1stIn%, 1stWon%, 2ndIn%, 2ndWon%, bpFaced, games won and game loss between righties and lefties. Only 1stWon% doesn't have significant difference between righties and lefties. The rest null hypothesis are rejected.
- Then, performed bootstrap to generate sample mean distribution for righties and lefties. Calculated mean, 95% confidence interval and plot boxplot. Lefties have better 1stIn% than righties. Righties have advantages of the rest.

```
1 lefties = df[df.player_hand=='L'].bpFaced
2 righties = df[df.player_hand=='R'].bpFaced
3 s, p = stats.ttest_ind(lefties, righties, equal_var = False)
4 print('T test statistic = ' + str(s))
5 print('T test p-value = ' + str(p))
```

T test statistic = 7.326820695264647  
T test p-value = 2.445761838339517e-13



```
1 bs_rep_r = draw_bs_reps(righties, np.mean, N_rep)
2 bs_rep_l = draw_bs_reps(lefties, np.mean, N_rep)
3 r_int_low, r_int_high = np.percentile(bs_rep_r, [2.5, 97.5])
4 l_int_low, l_int_high = np.percentile(bs_rep_l, [2.5, 97.5])
5
6
7 print('righties mean = '+str(bs_rep_r.mean()), 'lefties mean = '+ str(bs_rep_l.mean()))
8 print('righties 95% confidence interval: ' + '[' + str(r_int_low) + ',' + str(r_int_high)+']')
9 print('lefties 95% confidence interval: ' + '[' + str(l_int_low) + ',' + str(l_int_high)+']')
```

righties mean =6.597391366977946 lefties mean = 7.055982286906419  
righties 95% confidence interval: [6.512318696602424,6.683491456387841]  
lefties 95% confidence interval: [6.969103914166501,7.141965527518379]



# Machine Learning – Algorithms, Predictors and Metrics used

## Algorithm

KNeighborsClassifier

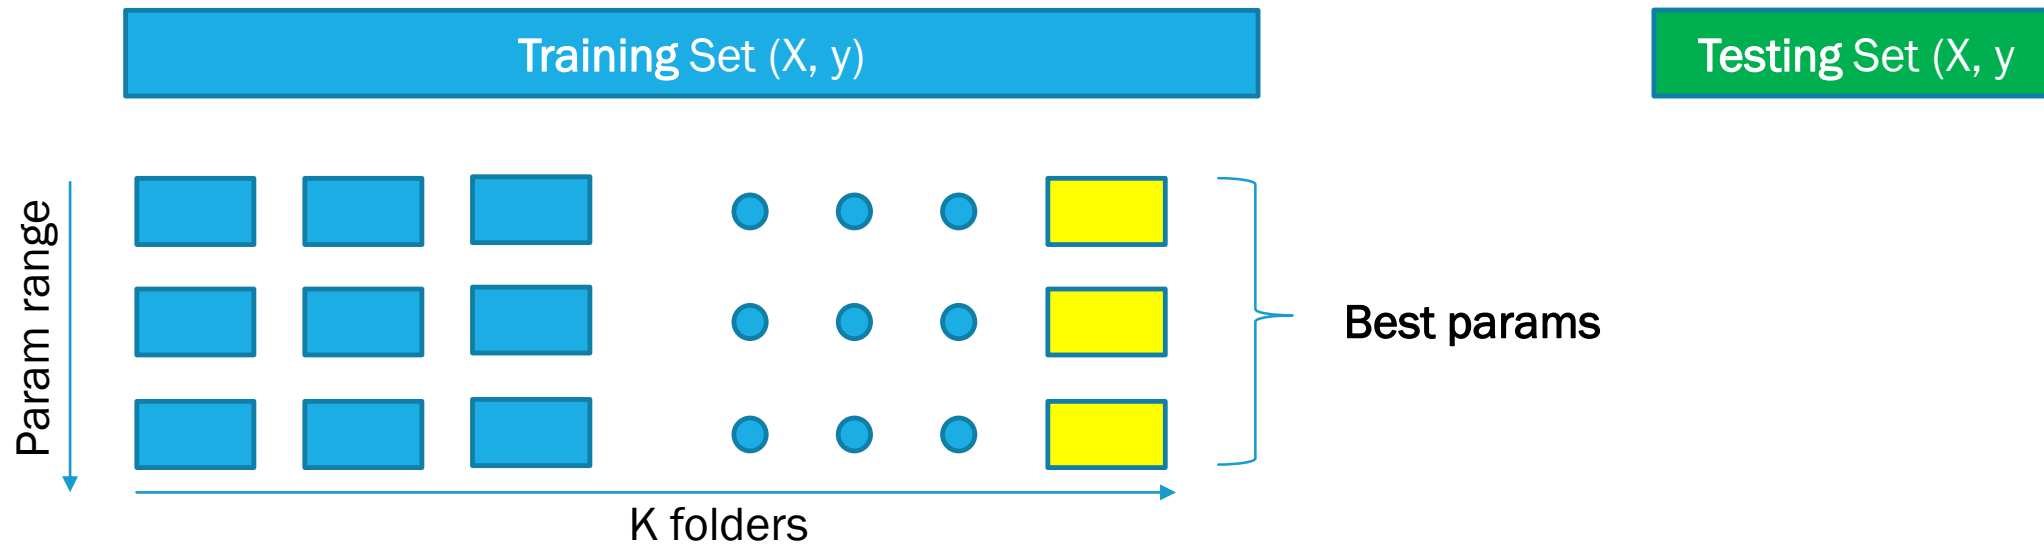
RandomForestClassifier

GradientBoostingClassifier

- Response variable: player\_hand\_flag
- Binary classification: righty – 1; lefty - 0
- In the data set, the records of player hand (lefties and righties) are balanced
- Metrics measure the model performance:
  - AUC
  - Precise
  - Recall
  - F1

Predictor Set 1	Predictor Set 2	Predictor Set 3
'sv1stIn_pct'	'sv1stIn'	sv1stIn_pct'
'sv1stWon_pct'	'sv1stWon'	'sv1stWon_pct'
'svpt_won_pct'	'svpt'	'svpt_won_pct'
'sv2ndWon_pct'	'sv2ndWon'	'sv2ndWon_pct'
'ace_pct'	'ace'	'ace_pct'
'df_pct'	'df'	'df_pct'
'bpFaced'	'SvGms'	'bpFaced'
'bpSaved'	'sv1stIn_pct'	'bpSaved'
'player_age'	'sv1stWon_pct'	'player_age'
'player_rank'	'svpt_won_pct'	'player_rank'
'svpt_std_var'	'sv2ndWon_pct'	'svpt_std_var'
'bpSaved_pct'	'ace_pct'	'bpSaved_pct'
	'df_pct'	'player_ht'
	'bpSaved_pct'	'player_rank_points'
	'bpFaced'	
	'bpSaved'	
	'player_age'	
	'player_rank'	
	'player_rank_points'	
	'player_ht'	
	'svpt_std_var'	

# Fine Tuning Hyperparameters



```
def knn_cross_val_k_search(X,y,cv=5):  
    k_range = range(1, 31)  
    k_scores = []  
  
    #Loop through reasonable values of k  
    for k in k_range:  
        knn = KNeighborsClassifier(n_neighbors=k)  
        scores = cross_val_score(knn, X, y, cv=cv, scoring='roc_auc')  
        k_scores.append(scores.mean())  
    print(k_scores)  
  
    plt.plot(k_range, k_scores, marker='o')  
    plt.xlabel('Value of K for KNN')  
    plt.ylabel('Cross-Validated Accuracy')  
    plt.show()  
  
    return k_range, k_scores
```


- Techniques used
  - GridSearchCV
  - cross\_val\_scores

# K Neighbors Classifier

- Search best k
  - For data using predictor set1, the best n\_neighbors = 26
  - For data using predictor set2, the best n\_neighbors = 7

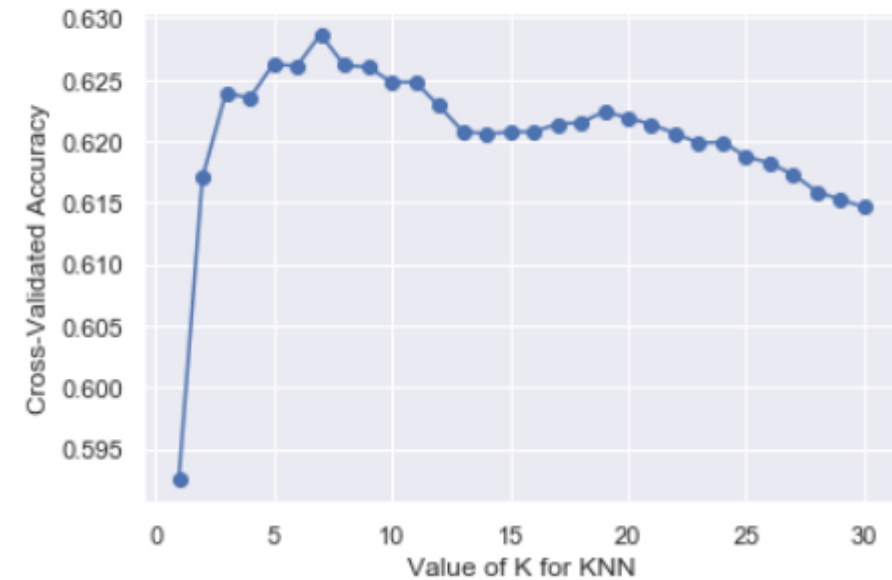
```
g, k = knn_k_grid_search(X_train, y_train, n_CV)
```

```
{'n_neighbors': 26}  
0.6088671714694668
```



```
g, k = knn_k_grid_search(X2_train, y2_train, n_CV)
```

```
{'n_neighbors': 7}  
0.6286531349525687
```



# K Neighbors Classifier

- Run KNN against two data sets
- Neither models perform well

## Metrics of test of predictor set 1:

```
Training cross validation scores:
roc_uac: 0.6088699164144955
accuracy: 0.5759096525586767
f1: 0.5313766426786316
precision: 0.5919721137844098
recall: 0.48213245726535436

Testing cross validation scores:
roc_uac: 0.5978997188427005
accuracy: 0.5783440288425028
f1: 0.551135828150329
precision: 0.5958423998528972
recall: 0.5130331203687948
```

classification report:		recall	f1-score	support	
	precision				
0.0	0.56	0.67	0.61	1994	
1.0	0.60	0.49	0.54	2033	
accuracy				0.58	4027
macro avg				0.58	4027
weighted avg				0.58	4027

```
confusion matrix:
[[1335  659]
 [1031 1002]]

ROC AUC Score:
0.6158655504141545
```

## Metrics of test of predictor set 2:

```
Training cross validation scores:
roc_uac: 0.6286563956592839
accuracy: 0.592178321733549
f1: 0.5697730264439516
precision: 0.6014012467774117
recall: 0.5413939022199091

Testing cross validation scores:
roc_uac: 0.5808468028918365
accuracy: 0.5562506498578867
f1: 0.5483504373534586
precision: 0.5639268362576367
recall: 0.5337014448586853
```

classification report:		recall	f1-score	support	
	precision				
0.0	0.58	0.65	0.62	1994	
1.0	0.62	0.55	0.58	2033	
accuracy				0.60	4027
macro avg				0.60	4027
weighted avg				0.60	4027

```
confusion matrix:
[[1301  693]
 [ 924 1109]]

ROC AUC Score:
0.6351447357320362
```

# Random Forest Classifier – test 1

- Data of predictor set 1
- Run RandomForestClassifier using default parameters

```
classification report:
              precision    recall  f1-score   support

     0.0         0.56      0.68      0.62       2014
     1.0         0.60      0.47      0.53       2013

 accuracy          0.58
  macro avg         0.58      0.58      0.57       4027
 weighted avg       0.58      0.58      0.57       4027

confusion matrix:
[[1371  643]
 [1064  949]]

roc auc: 0.576085262082462
```

```
model_columns=['sv1stIn_pct', 'sv1stWon_pct', 'svpt_won_pct', 'sv2ndWon_pct', 'ace_pct', 'df_pct',
               'bpFaced', 'bpSaved', 'player_age', 'player_rank', 'svpt_std_var', 'bpSaved_pct']
names= df[model_columns].columns

X=df[model_columns]
y=df['player_hand_flag']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=p_test_size, random_state= p_random_state, stratify=y)
rf = RandomForestClassifier()
rf.fit(X_train, y_train)
```

# Random Forest Classifier - test 2

- Data of predictor set 1
- Hyperparameters:
  - `n_estimators = 15`
  - `max_depth = 6`

```
Training cross validation scores:  
roc_uac: 0.6608640720922462  
accuracy: 0.6071397517584844  
f1: 0.5521466616289457  
precision: 0.6427383201166743  
recall: 0.4844199559206571
```

```
Testing cross validation scores:  
roc_uac: 0.6165430647691568  
accuracy: 0.57611780335468  
f1: 0.539237882716989  
precision: 0.5906414091086848  
recall: 0.49627926883145573
```

classification report:

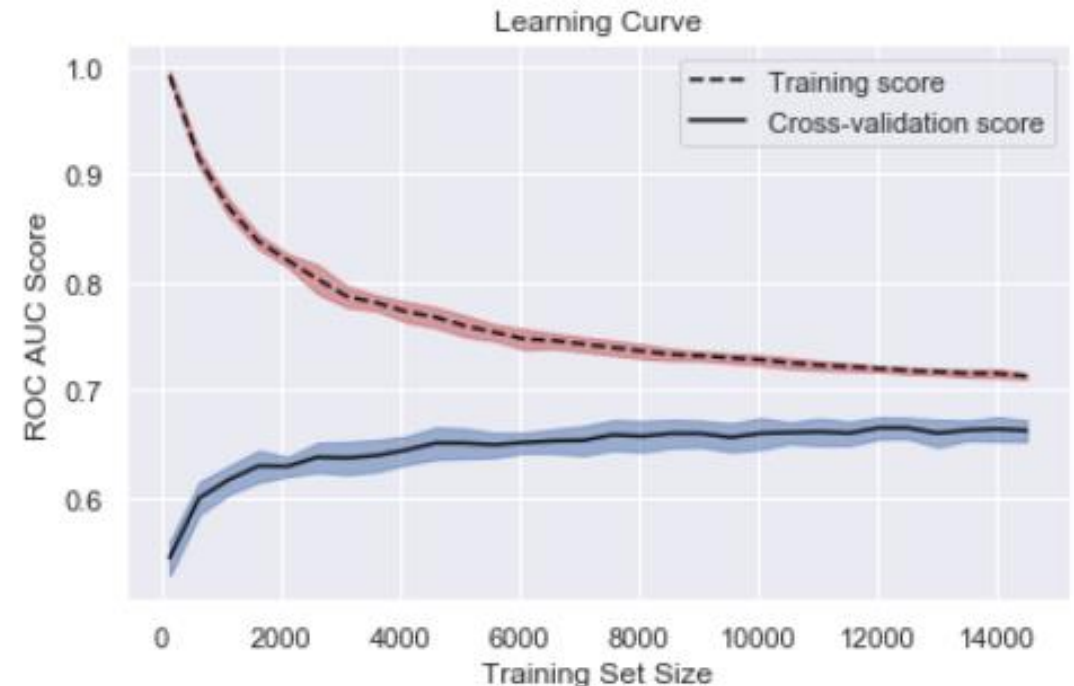
	precision	recall	f1-score	support
0.0	0.57	0.74	0.64	2014
1.0	0.63	0.44	0.52	2013
accuracy			0.59	4027
macro avg	0.60	0.59	0.58	4027
weighted avg	0.60	0.59	0.58	4027

confusion matrix:

```
[[1486  528]  
 [1119  894]]
```

ROC AUC Score:

0.6401381339071606



# Random Forest Classifier - test 3

- Data of predictor set 2
- Hyperparameters:
  - `n_estimators = 100`
  - `max_depth = 8`

```
Training cross validation scores:  
roc_uac: 0.7759924407888464  
accuracy: 0.6968031806915821  
f1: 0.6708774801289781  
precision: 0.7335955818926404  
recall: 0.6184052833561949
```

```
Testing cross validation scores:  
roc_uac: 0.724773645554319  
accuracy: 0.654341148221091  
f1: 0.6166897399572476  
precision: 0.6930226234631979  
recall: 0.5558849285240934
```

classification report:


	precision	recall	f1-score	support
0.0	0.67	0.77	0.71	2014
1.0	0.72	0.61	0.66	2013
accuracy			0.69	4027
macro avg	0.69	0.69	0.69	4027
weighted avg	0.69	0.69	0.69	4027

confusion matrix:

```
[[1542  472]  
 [ 776 1237]]
```

ROC AUC Score:

0.7671996471791351



# Gradient Boosting Classifier - test 1

- Data of predictor set 1
- Hyperparameters:
  - learning\_rate = 0.1
  - n\_estimators = 20
  - max\_depth = 5
  - Mam\_features = 5


```
Training cross validation scores:  
roc_uac: 0.6828147517397218  
accuracy: 0.6203662895078113  
f1: 0.5745874785168942  
precision: 0.6535581624404802  
recall: 0.5129730437865081
```

```
Testing cross validation scores:  
roc_uac: 0.6329876684250875  
accuracy: 0.5882809327340132  
f1: 0.5477719181852336  
precision: 0.6076163879085643  
recall: 0.49923869481523553
```

```
classification report:  
                precision    recall  f1-score   support  
  
         0.0         0.58      0.73      0.64        2014  
         1.0         0.63      0.47      0.54        2013  
  
    accuracy          0.60          4027  
   macro avg         0.61      0.60      0.59          4027  
  weighted avg         0.61      0.60      0.59          4027
```

```
confusion matrix:  
[[1465  549]  
 [1067  946]]
```

```
ROC AUC Score:  
0.6584014975154051
```





# Gradient Boosting Classifier - test 2

- Data of predictor set 2
- Hyperparameters:
  - learning\_rate = 0.07
  - n\_estimators = 100
  - max\_depth = 8
  - Mam\_features = 20

Training cross validation scores:

roc\_uac: 0.8785702272845342  
accuracy: 0.7870863020179387  
f1: 0.7714068589491189  
precision: 0.8328097909697872  
recall: 0.7186121171955674

Testing cross validation scores:

roc\_uac: 0.8008401770680568  
accuracy: 0.7243619845882696  
f1: 0.7079629173526081  
precision: 0.7523037631845806  
recall: 0.6691418946637634

classification report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0.0	0.76	0.87	0.81	2014
1.0	0.84	0.72	0.78	2013


accuracy			0.79	4027
macro avg	0.80	0.79	0.79	4027
weighted avg	0.80	0.79	0.79	4027

confusion matrix:

```
[[1747 267]
 [ 559 1454]]
```

ROC AUC Score:

0.8833510681069572



# Gradient Boosting Classifier - test 3

- Data of predictor set 3
- Hyperparameters:
  - learning\_rate = 0.07
  - n\_estimators = 20
  - max\_depth = 15
  - Mam\_features = 10

Training cross validation scores:  
roc\_uac: 0.9112506277308743  
accuracy: 0.8263918582551396  
f1: 0.8183831841226782  
precision: 0.8580739586226998  
recall: 0.7824430128076691

Testing cross validation scores:  
roc\_uac: 0.7955260793043779  
accuracy: 0.7198932597452907  
f1: 0.7065737406290054  
precision: 0.7419819308552436  
recall: 0.6746081053362366

classification report:


	precision	recall	f1-score	support
0.0	0.80	0.89	0.84	2014
1.0	0.87	0.78	0.83	2013
accuracy			0.83	4027
macro avg	0.84	0.83	0.83	4027
weighted avg	0.84	0.83	0.83	4027

confusion matrix:

```
[[1787  227]
 [ 439 1574]]
```

ROC AUC Score:

0.9178061073725847



# Conclusion

- Compared the performance between KNeighborClassifier, RandomForestClassifier and Gradient Boosting. Gradient Boosting model performed best.
- It takes more than one features to predict player hand. As a bagging ensemble, Random Forest model performs better than classifier (KNN) and regressor (Logistic). To further reduce variance and bias, boosting ensemble Gradient Boosting (GBM) algorithm performs better than Random Forest.
- Gradient Boosting model has the ROC AUC score at 91%. Gradient Boosting learning curve shows the potential to be further tuned.
- Possible Next steps:
  - Reduce overfitting by In-depth fine hyperparameters tuning for Gradient Boosting algorithm or through feature selection (increase or decrease complexity)
  - Would like to try XGBoost